

Augmenting Romer's Endogenous Growth Theory for the Age of Cognitive Automation

Jean Louis Van Belle, *Drs, MAEc, BAEC, BPhil* & Google Gemini, 4-5 October 2026

Abstract

Paul Romer's (1990) foundational endogenous growth model treats knowledge (A) as a non-rival, partially excludable asset driven exclusively by the allocation of human research labor (L_A). While robust for the early software and industrial eras, the Romer framework carries a critical blind spot in the era of Artificial Intelligence: it assumes technology operates purely as a complement to human intellect, failing to account for direct cognitive factor substitution.

This theoretical note bisects human knowledge labor into distinct senior (L_{senior}) and junior (L_{junior}) cohorts, and introduces automated cognitive capital (K_{AI}) via a Constant Elasticity of Substitution (CES) knowledge production function. Crucially, we formalize L_{senior} as an endogenous accumulation function fed directly by historical junior upskilling. We mathematically demonstrate how contemporary asymmetric fiscal signals distort corporate optimization—inducing a rational short-term freeze on junior hiring that leads to long-term institutional capability bankruptcy through exponential skills debt.

Finally, we move beyond pure theory and modeling to address the real-world constraints of geopolitical arbitrage and regulatory leakage. We evaluate the international legal vulnerabilities of imposing external digital tariffs under the 1998 WTO Moratorium and GATS Article XVII. In doing so, **we propose specific, actionable steps national policy-makers can take immediately to correct factor-price distortions while staying completely within international rules on tariffs and taxation.**

Keywords: Artificial Intelligence; Corporate Strategy; Human Capital Atrophy; AI Governance; Endogenous Growth Theory; Managing GenAI; Tax Neutrality; WTO Compliance; GATS.

Contents

Abstract	0
1. The Classical Baseline: Romer (1990).....	2
2. The AI-Augmented Framework: Factoring Cognitive Capital	2
3. The Dynamic Endorsement of Senior Capital	3
4. The Microeconomic Optimization Trap	3
5. The System Black Box Threshold and Macro Stagnation.....	3
6. Policy Intervention: Adjusting the Production Parameters	5
7. Geopolitical Arbitrage and WTO Compatibility of Digital Tariffs.....	5
7.1 The Regulatory Leakage Paradox	5
7.2 The Primary Obstacle: The WTO Electronic Transmission Moratorium.....	6
7.3 The Superiority of the Revenue-Neutral Internal Shift	6
7.4 GATS Compliance Comparison Matrix	7

Annex A: The Microeconomic Corporate AI Growth Model	8
Annex B: Introducing Black-Swan Events	9
Concept.....	9
Real-World Historical Examples	9
Similarities between the System Black Box and Black-Swan Modelling	10
Annex C: AI and the Synthesis of Value Theories	11
The Generational and Ideological Dissolution of Value	11
How AI Breaks the Traditional Binary	11
1. The Crisis of the Labor Theory (The Socialist Heart)	11
2. The Crisis of the Subjective Market Theory (The Liberal Mind)	11
3. The Valuation Paradox: Collusion vs. Destructive Competition	11
Factor-Price Neutrality: the Liberal Mind Serving the Socialist Heart	12
The Operational Duck Test: Sidestepping Metaphysical Traps	13
Annex D: Microeconomic Calibration and the Belgian Macro-Fiscal Case Study	15
1. The Sovereignty Paradox: The US-China Infrastructure Duopoly	15
2. The Downstream Solution: The Reverse-Charge Token Mechanism	15
3. The Compliance Lever: Corporate OpEx Deductibility.....	16
4. Neutralizing the Open-Source and Local Compute Loophole	17
5. The 1:1000 Microeconomic Micro-Foundations	17
6. The Belgian Macro-Fiscal Extrapolation	18
The Math of the Failed Ad Valorem Model	19
The Volumetric Surcharge Alternative.....	19
Annex E: Future Refinements of the Model	21
1. Strengths Across Both Papers	21
2. Targeted Improvements for the <i>Corporate Playbook</i> (companion paper)	21
A. Formalize the Dimension of System Complexity Ct	21
B. Expand "Verification Density" into an Audit Formula	21
3. Targeted Improvements for the <i>Augmented Romer Model</i> (this paper)	22
A. Explicit Feedback from Knowledge Stock A to Senior Atrophy γ	22
B. Address the Frontier Model Pricing Trend	22
C. The Open-Source GPU Proxy Enforcement	22
Summary Assessment	22
Annex F: The Micro-Financial Mechanics of Cognitive Capital.....	23
1. The Bertrand-Schumpeter Valuation Paradox	23
Formalizing the Monopoly Fallacy vs. Marginal Cost Pricing	23
2. Disaggregating the Service Flow: The Three-Tier Token Metric	24
3. Closing the "Shadow AI" / Private Account Tax Loophole	24
References.....	27

1. The Classical Baseline: Romer (1990)

In the standard Romer framework, the aggregate production function of an economy is expressed as:

$$Y = K^\alpha (A \cdot L_Y)^{1-\alpha}$$

Where Y is total output, K is physical capital, L_Y is labor allocated to final goods production, and A is the stock of knowledge/ideas. The accumulation of knowledge is driven by an independent non-rival production function:

$$\dot{A} = \delta \cdot A \cdot L_A$$

Where L_A represents the total allocation of human research labor (engineers, scientists, analysts) and δ is a productivity parameter. Romer implicitly relies on two axioms that modern AI fundamentally violates:

1. **The Non-Rivalry of Aggregate Ideas:** Ideas exist as abstract public/club goods that scale linearly across the macroeconomy.
2. **Exogenous Labor Inputs:** The capacity of L_A to direct and generate ideas is an exogenous pool of human talent.

2. The AI-Augmented Framework: Factoring Cognitive Capital

In this paper, we shift the analytical focus from macro-communal ideas to **firm-specific proprietary cognitive capital**. When a firm deploys an array of advanced generative AI models (K_{AI}), it does not merely tap into a public Romerian idea; it builds context-dependent, heavily excludable, localized corporate workflows. Furthermore, AI acts uniquely as **capital behaving like intellectual labor**. To model this factor substitution, we split the traditional human research labor pool (L_A) into two structurally dependent cohorts:

- $L_{senior}(t)$: Veteran human experts possessing deep conceptual, architectural, and diagnostic mental models.
- $L_{junior}(t)$: Entry-level apprentices historically tasked with lower-level execution (drafting, boilerplate engineering, basic synthesis).

We reformulate the aggregate knowledge generation function of the enterprise using a **Constant Elasticity of Substitution (CES) production function**:

$$\dot{A} = \psi \cdot A^\phi \left[\beta (L_{senior} \cdot L_{junior})^\rho + (1 - \beta) K_{AI}^\rho \right]^{\frac{1}{\rho}}$$

Where ψ is an efficiency coefficient, $\beta \in [0,1]$ represents the distributional weight of human-collaborative labor, and $\rho \leq 1$ dictates the elasticity of substitution between human apprenticeships and automated AI agents ($\sigma = \frac{1}{1-\rho}$).

In the modern operational landscape, for entry-level execution and baseline drafting tasks, the elasticity of substitution approaches infinity ($\rho \rightarrow 1$), meaning K_{AI} and the junior human workforce operate as near-perfect substitutes.

3. The Dynamic Endorsement of Senior Capital

The fundamental structural flaw of modern corporate planning lies in treating L_{senior} as an exogenous, static variable. In reality, senior human capital is a **dynamic, endogenous accumulation function** dependent entirely on the historical pipeline of junior training.

We formalize the law of motion for senior expertise as:

$$L_{senior}(t) = L_{senior}(0)e^{-\gamma t} + \int_0^t \theta \cdot L_{junior}(\tau - \lambda) d\tau$$

Where:

- γ is the organic industry turnover and retirement rate.
- θ is the conversion efficiency coefficient of a junior professional transitioning into a senior practitioner.
- λ represents the structural time lag of human cognitive apprenticeship (typically $3 \leq \lambda \leq 5$ years).

4. The Microeconomic Optimization Trap

A cost-minimizing corporate manager seeks to optimize the input factor mix based on relative factor prices. Let the total operational cost of the knowledge infrastructure be:

$$C_{ops} = w_s L_{senior} + (w_j + \tau_p) L_{junior} + p_k K_{AI}$$

Where w_s and w_j are human wages, τ_p represents payroll taxes and social insurance surcharges, and p_k is the market price of automated API tokens.

Because contemporary tax code structures heavily tax human employment via τ_p while permitting p_k to be discharged cleanly as an operating expense (OpEx) with zero long-term social liability, the pricing signal is severely distorted:

$$\frac{w_j + \tau_p}{p_k} \rightarrow \infty \quad \text{as} \quad p_k \rightarrow 0$$

Faced with this massive asymmetry, a rational manager executing short-term optimization will set entry-level human recruitment to zero:

$$L_{junior}(t) = 0 \quad \forall \quad t \geq 0$$

5. The System Black Box Threshold and Macro Stagnation

When L_{junior} is suppressed to zero, the dynamic accumulation function derived in Section 3 loses its influx vector, degenerating into pure exponential decay:

$$L_{senior}(t) = L_{senior}(0)e^{-\gamma t}$$

Substituting this back into our knowledge production function reveals a terrifying macro-economic lag. For the initial horizon $t \in [0, \lambda)$, corporate output ($\dot{A}$) appears hyper-efficient;


costs have plummeted while legacy senior experts comfortably leverage K_{AI} to manage complex operations.

However, as $t \rightarrow \lambda$, the system hits a critical threshold (t^*) and collapses into an irreversible skills debt:

$$t^* = \frac{\ln(L_{senior}(0)/K_{AI}(0))}{\gamma + r}$$

Beyond t^* , the firm's aggregate capacity to establish strategic constraints, spot model hallucinations, or debug complex architectural failures drops below the complexity threshold of the software infrastructure itself. The enterprise loses endogenous capacity, entering a permanent state of structural stagnation.

This is illustrated by the intersection of the senior labor and AI complexity curves in Scenario A below (unmanaged growth).



6. Policy Intervention: Adjusting the Production Parameters

As we can see, the rational short-term micro-optimization of production factors by organizations adopting AI, leads to systems failure in a timeframe that may be as short as four to five years. This is, obviously, an undesirable longer-term outcome. To correct this, market incentives should be changed to incorporate longer-term consideration. The theory of public economics tells us that, when market incentives are distorted, fiscal architecture can be used to directly target the real factor price ratio. By implementing a **revenue-neutral tax shift** that places a minor surcharge on high-volume autonomous API tokens ($p_k + \tau_{AI}$) and uses 100% of the proceeds to diminish human payroll levies ($\tau_p \rightarrow 0$), the relative factor cost becomes more balanced:

$$\frac{w_j + 0}{p_k + \tau_{AI}}$$


This structural recalibration changes the corporate optimization math without increasing the total tax burden on enterprises. It mathematically establishes the *Human-in-the-Loop* augmented configuration as the most competitive factor mix, preserving both immediate productivity velocity and the long-term human talent runway.

7. Geopolitical Arbitrage and WTO Compatibility of Digital Tariffs

7.1 The Regulatory Leakage Paradox

The revenue-neutral tax shift proposed in Section 6 relies on the sovereign state's capacity to adjust factor price mix incentives within its domestic jurisdiction. However, in a hyper-globalized, frictionless digital economy, cognitive capital is highly mobile. If a single economic bloc (e.g., the European Union) unilaterally implements an operational automation levy on high-volume autonomous API tokens ($p_k + \tau_{AI}$), it introduces a severe risk of geopolitical arbitrage.

Indeed, because software code and algorithmic instructions are frictionless, an enterprise can easily exploit jurisdictional boundaries:



To counter this arbitrage, policy thinkers frequently propose a digital variant of the **Carbon Border Adjustment Mechanism (CBAM)**—applying an *Algorithmic Labor Displacement Levy* to imported digital services generated entirely by untaxed overseas autonomous agents. However, executing this mechanism requires navigating a minefield of international trade treaties.

7.2 The Primary Obstacle: The WTO Electronic Transmission Moratorium

The primary international legal barrier to any cross-border algorithmic levy is the **World Trade Organization (WTO) Moratorium on Customs Duties on Electronic Transmissions**.

Established in 1998, this multilateral agreement explicitly prohibits member states from imposing tariffs or customs duties on any digital goods or services delivered electronically. Because an overseas AI agent delivers its output (codebases, legal briefs, market analyses) exclusively via electronic data transmissions, any direct tariff applied at the digital border constitutes an explicit violation of the active WTO moratorium. Unilaterally breaking this commitment risks dismantling global digital trade predictability, inviting immediate retaliatory trade sanctions from global trading partners.

The Legal Resolution: GATS Article XVII and National Treatment

To prevent an AI border adjustment from degenerating into a tool for **unilateral protectionism**—and to insulate it from being struck down by trade tribunals—it cannot legally operate as a *customs duty applied at the border*. Instead, it must be structured as an **internal non-discriminatory consumption tax** governed strictly by the **General Agreement on Trade in Services (GATS)**. Indeed, to survive a challenge under GATS, the fiscal framework must flawlessly respect the foundational principle of **National Treatment (GATS Article XVII)**:

$$\text{Tax Rate } (\tau_{AI}) \text{ on Imported AI Token} \equiv \text{Tax Rate } (\tau_{AI}) \text{ on Domestic AI Token}$$

The regulatory levy cannot discriminate based on geographic origin:

1. If a domestic enterprise imports automated cognitive output from a foreign frontier model provider, it faces the exact same internal token fee as a firm utilizing a locally hosted, domestic AI infrastructure.
2. **The tax must target the *nature of the production factor substitution itself*** (capital replacing human entry-level positions), rather than the nationality of the technology provider.

7.3 The Superiority of the Revenue-Neutral Internal Shift

This legal architecture demonstrates why the external "digital border tariff" is a fragile policy mechanism, easily abused by protectionist factions to shelter uncompetitive domestic industries. The most robust, legally insulated defense against the factor-extinction horizon remains the **unilateral adjustment of domestic factor floor prices**. By utilizing tax revenue on automated capital to significantly lower payroll taxes and social security contributions on human workers ($\tau_p \rightarrow 0$), the sovereign state changes the optimization ratio without ever touching a trade boundary. It respects the GATS National Treatment principle entirely, avoids the WTO electronic moratorium trap, and alters the relative price signal to structurally reinforce the competitive position of the *Human-in-the-Loop* augmented professional.


7.4 GATS Compliance Comparison Matrix

The following table offers an overview of the legal boundaries and why and how our proposal overcomes the fragility of imposing external tariffs on API consumption and – importantly – allows national policy-makers to move ahead immediately.

Proposed Fiscal Instrument	WTO Moratorium Compatibility (1998 Electronic Trade)	GATS Article XVII Compatibility (National Treatment)	Retaliatory Trade War Risk	Primary Structural Utility
Unilateral Digital Border Levy (<i>Tariff on Foreign APIs</i>)	✗ Direct Violation (Bans duties on electronic transmissions)	✗ Direct Violation (Explicitly discriminates by origin)	● Extreme (Triggers immediate tech sector counters)	None; highly fragile and prone to immediate WTO invalidation.
Global Algorithmic CBAM (<i>Tariff on AI Output Services</i>)	✗ High Risk (Struck down if service is purely digital)	⚠ Conditional (Only legal if local models pay identically)	● High (Perceived as targeted digital protectionism)	Safeguards borders but creates immense geopolitical friction.
Internal Revenue-Neutral Tax Shift (<i>Our Framework</i>)	100% Compliant (Bypasses border; operates purely internally)	100% Compliant (Taxes <i>factor substitution</i> , not provider origin)	● Zero Risk (Pure exercise of domestic fiscal sovereignty)	Maximum; structurally shifts local input prices without touching trade laws.

Annex A: The Microeconomic Corporate AI Growth Model

We created a basic model engine allowing the readers of this paper to verify our model and extend it. This engine is available as Model1.py on our GitHub site: <https://github.com/jeanlouisvanbelle/AI-Economics>. It incorporates the basic equations and generates the visual which we used in the main body of this paper:



Even if this model is quite simple (no feedback loops), it sufficiently illustrates two completely divergent physical outcomes:

1. **The Non-Equilibrium Zone (Scenario A):** Because L_{junior} is dropped to zero, the inflow of new human capability stops. Senior capital decays unreplenished. Meanwhile, unvetted code and automated tasks cause technical debt and complexity to skyrocket at a compound rate of 12% per year. The functions cross in Year 4 after the freeze on the hiring on juniors, diving into a state of structural bankruptcy where the company can no longer audit its own systems.
2. **The Dynamic Equilibrium Track (Scenario B):** When your playbook rules are enforced, the firm injects a steady stream of junior apprentices who are trained utilizing AI-human training models as described in our [Corporate Playbook for AI-Augmented Apprenticeships](#) (Van Belle, 2026) such as *Adversarial Audits* and *Sandboxed Iterations*. Even though there is a strict 3-year cognitive time lag ($\lambda = 3$) while they learn, by year 4 the trained apprentices cleanly offset senior retirements. The system establishes a **Dynamic Steady State**, where senior capacity $L_{senior}(t)$ stabilizes and starts leveling off above the controlled, refactored complexity curve.

Annex B: Introducing Black-Swan Events

Concept

A **black-swan event** is a catastrophic, completely unexpected systemic failure that occurs when an artificial intelligence system or automated algorithm encounters real-world conditions it was never exposed to during its training phase.

The term combines two foundational concepts from economics and technology:

1. **The Black Swan** ([Nassim Nicholas Taleb, 2007](#)): An event that is an extreme outlier, carries an immense, widespread impact, and is rationalized only *after* it happens (e.g., the 2008 financial crash or the 2020 global pandemic).
2. **Algorithmic Model Drift**: The mathematical reality that AI models are static snapshots trained on historical data. When the real world abruptly changes, the model's logic breaks down entirely because it can only predict the future based on the past.

When these two realities collide, the AI model does not just make a minor error; it enters a state of **silent, hyper-accelerated failure**, generating highly confident but completely wrong or dangerous outputs.

Real-World Historical Examples

To ground this mathematically, we do not have to look to sci-fi; we have stark historical precedents where unvetted automation triggered immediate, multi-billion-dollar market collapses:

1. The Knight Capital Group Collapse (2012)

- **The Error**: A software update left a piece of legacy, unvetted algorithmic trading code active in production. When the market opened, the autonomous system entered an uncontrolled loop, buying and selling millions of shares in an erratic pattern.
- **The Impact**: Because the firm lacked sufficient human-in-the-loop oversight to instantly understand the automated breakdown, **the algorithm lost \$440 million in exactly 45 minutes**, bankrupting the entire corporation before lunch.


2. The Zillow "Offers" Property Pricing Disaster (2021)

- **The Error**: Real estate platform Zillow deployed an algorithmic pricing engine to autonomously buy houses based on historical market data. When the post-pandemic supply chain disrupted the labor and material markets, housing costs shifted abnormally—creating a classic "black swan" data shift.
- **The Impact**: The algorithm failed to understand the anomaly and kept aggressively over-purchasing thousands of homes at inflated prices. Zillow had to shut down the entire division, lay off 25% of its workforce, and **write off over \$500 million in losses** because humans trusted the machine blindly.

Similarities between the System Black Box and Black-Swan Modelling

In this paper, we showed what happens if a company freezes junior hiring: they enjoy high profit margins during the initial window because their veteran seniors can easily manage routine AI outputs. However, if the firm hits Year 5, the model tells us it is likely to encounter an unprecedented real-world crisis. The question is not *if*, but *when* it will encounter such crisis. We think it might be triggered by black-swan events such as the ones described above. In such cases, an event combines with two organizational *states* the organization has maneuvered itself into because of its short-term optimization of its labor force. We illustrate this below.

THE CATASTROPHIC INTERSECTION BEYOND YEAR 4.39



Annex C: AI and the Synthesis of Value Theories

The Generational and Ideological Dissolution of Value

For over two centuries, political economy has been defined by the binary tension between the **Labor Theory of Value** (the bedrock of socialist critique) and the **Subjective/Marginal Theory of Value** (the foundation of liberal market economics).

In the lifecycle of political philosophy, it is often observed that the young 'heart' defaults to the left—an intuitive focus on the collective good and the intrinsic value of human labor. However, when the adult observes and reflects on real-world experience with the inefficiencies of state execution, it typically gives rise to a more liberal mindset—a healthy skepticism of bureaucratic intervention, regulation, and top-down distortions.

The advent of automated cognitive capital (K_{AI}) throws both frameworks into structural disarray, exposing them not as eternal laws, but as historical artifacts of an era when cognitive processing was inextricably bound to biological human brains.

How AI Breaks the Traditional Binary

1. The Crisis of the Labor Theory (The Socialist Heart)

The socialist axiom asserts that economic value is fundamentally crystallized human labor. Capital is merely "dead labor" used to enhance or exploit living workers.

When a generative AI engine processes millions of legal documents or drafts optimized codebases for a marginal cost approaching zero ($p_k \rightarrow 0$), it generates massive, actionable intellectual output with near-zero active human labor time. If value is strictly human labor, this hyper-abundant cognitive output should possess no economic reality. Yet, it actively alters corporate productivity and displaces workers. The "socialist heart" must confront an economy where high-level cognition has been permanently decoupled from human meatware.

2. The Crisis of the Subjective Market Theory (The Liberal Mind)


Concurrently, AI breaks classical market optimization. Liberal economics relies on *scarcity* to allocate resources via price signals. By making knowledge hyper-abundant and stripping away the marginal cost of execution, the market mechanism paradoxically devalues the core asset of the "knowledge society." Even if a consumer derives immense subjective utility from an AI solution, cut-throat competition among model providers and the zero-marginal-cost boundary force the market price down toward zero.

3. The Valuation Paradox: Collusion vs. Destructive Competition

This internal breaking of market mechanics explains the deeply abnormal financial environment surrounding modern AI technology companies.

THE AI VALUATION PARADOX

The Fragmented Cycle of Speculative Capital vs. Collapsing Product Margins



Popular critique frequently mischaracterizes the dominance of large technology conglomerates as a symptom of cozy collusion or classic monopoly. The empirical reality is the exact opposite: it is an environment of **destructive, cut-throat competition**.

Standard economic valuation models (such as Discounted Cash Flow) completely fail to justify current AI stock prices based on current earnings. Investors clearly understand the staggering, civilization-shifting potential of cognitive automation, yet the market is currently failing to establish a stable mechanism to balance fair corporate rewards with sustainable competition.

Because the underlying product (the API token) faces a marginal cost structure collapsing toward zero, these enterprises cannot easily convert immense societal utility into stable, profitable cash flows. Consequently, the entire AI infrastructure thrives artificially on massive debt, speculative venture capital, and cross-subsidization from legacy business models, chasing highly uncertain future returns. Cut-throat competition is not a sign of a healthy, self-correcting market; it is a symptom of a structural malfunction where value cannot be properly captured or retained.

Factor-Price Neutrality: the Liberal Mind Serving the Socialist Heart

The policy intervention proposed in this paper—the internal, revenue-neutral tax shift—represents a pragmatic ideological synthesis that respects both the liberal mind and the socialist heart:

- **The Liberal Mind (The Mechanism):** It explicitly rejects heavy-handed regulatory bans, state-mandated pricing, or clumsy digital border tariffs that invite bureaucratic failure, protectionism, and international trade wars. It also accepts that corporate managers must operate as rational cost-minimizers.

- **The Socialist Heart (The Objective):** It insists that the ultimate purpose of an economic architecture is the preservation and cultivation of human capability and societal resilience. It recognizes that an economic system that allows its foundational human talent pipeline to undergo exponential decay ($H_0 e^{-\gamma t}$) for short-term balance sheet optimization is structurally bankrupt.

Indeed, the revenue-neutral shift ($\tau_{AI} \rightarrow \tau_p \rightarrow 0$) does not use state power to dictate subjective value or pick industrial winners.¹ Instead, it uses a clean, rule-based fiscal adjustment to correct a profound systemic distortion: **the fact that current tax codes heavily penalize biological human employment while heavily subsidizing algorithmic server transactions.**

By significantly shifting taxes in a way that is – overall – budget-neutral, the liberal mechanism of free corporate optimization is fully preserved while, at the same time, the microeconomic optimization math is refactored so that the market naturally selects the "Human-in-the-Loop" configuration as the most stable, competitive factor mix. In this manner, market efficiency (the mind) is structurally reconciled with the preservation of human capital and societal capability (the heart).

The Operational Duck Test: Sidestepping Metaphysical Traps


A common vulnerability in contemporary AI governance literature is the reliance on deep metaphysical or legalistic definitions of machine consciousness, intentionality, or human-equivalence to justify intervention. Proponents of "AI alignment" or "robot taxation" frequently get bogged down in proving whether a digital architecture truly "understands" its output or possesses independent moral agency.

This paper's framework explicitly rejects the necessity of resolving these philosophical debates. Instead, the macro-fiscal justification for factor-price neutrality relies on a highly pragmatic variation of the classical empirical duck test:

If something looks like a duck, swims like a duck, and quacks like a duck, then it may not be a duck, but it should be taxed like a duck.

¹ Throughout the paper, and also in this Annex, the functional shorthand notation $\tau_{AI} \rightarrow \tau_p \rightarrow 0$ denotes a structural fiscal transfer vector rather than an absolute mathematical limit. It signifies that revenues generated by the automation surcharge (τ_{AI}) are recycled to systematically drive down human payroll liabilities (τ_p). In macro-fiscal application, this does not mandate an immediate drop to absolute zero, but rather an asymptotic reduction toward a minimized factor floor price, removing the structural fiscal penalty on human employment.

THE OPERATIONAL DUCK TEST IN PUBLIC FINANCE



By decoupling the fiscal mechanism from the internal nature of the intelligence (whether biological meatware or algorithmic software), the proposed policy intervention becomes academically unassailable. We do not need to legally or spiritually define what the AI *is*; we merely observe what it *does* to corporate resource optimization. If automated cognitive capital behaves as a near-perfect substitute for entry-level human execution ($\rho \rightarrow 1$), its structural footprint on institutional capability decay ($H_0 e^{-\gamma t}$) mandates a balanced fiscal response based entirely on its operational realities rather than its metaphysical status.

Annex D: Microeconomic Calibration and the Belgian Macro-Fiscal Case Study

1. The Sovereignty Paradox: The US-China Infrastructure Duopoly

A critical challenge to any upstream fiscal intervention at the "basis" of technology production is the highly centralized reality of global computing infrastructure. The generative AI landscape is overwhelmingly structured as an asymmetric duopoly:

- **The United States** anchors the dominant hyperscale cloud infrastructure and frontier foundation models via **Amazon Web Services (AWS), Microsoft Azure, Google Cloud, and Oracle**. [1, 2]
- **China** commands the secondary sphere of influence, concentrated through **Alibaba Cloud, Huawei Cloud, and Tencent Cloud**. [1]
- **Europe and Regional Blocs** retain a minimal market share (~15%), represented by specialized local providers like **OVHcloud (France) or Hetzner (Germany)**, which remain largely dependent on foreign silicon infrastructure. [1, 2, 3]

If a small or mid-sized sovereign nation (e.g., Belgium) attempts to mandate that an overseas hyperscaler alter its core global billing architecture to collect a token-metered surcharge at the upstream source, it encounters a wall of jurisdictional friction. Unilateral extraterritorial enforcement on foreign corporate monopolies is practically impossible without deep bilateral cooperation or multilateral treaties with Washington or Beijing—prospects that are politically distant and highly volatile. [1, 2, 3]

2. The Downstream Solution: The Reverse-Charge Token Mechanism

To remain completely viable and implementable by individual national tax authorities without external dependency, the fiscal architecture must pivot **downstream**. Rather than attempting to tax the foreign *supplier* at the point of origin, the sovereign state levies the operational automation fee directly on the domestic *purchaser* at the point of consumption.

This mechanism directly borrows the established, internationally compliant framework of the **Destination Principle** used for cross-border digital services and business-to-business (B2B) Software-as-a-Service (SaaS) taxation. [1, 2, 3]

UNILATERAL DOWNSTREAM ENFORCEMENT FRAMEWORK



Under this downstream framework, when a local enterprise imports automated cognitive outputs from a foreign frontier model provider via API, the transaction triggers a mandatory **Reverse-Charge Token Assessment**:

1. The foreign provider issues a net invoice with zero localized tax, completely bypassing international trade blockages or the WTO Electronic Transmission Moratorium.
2. The domestic enterprise is legally required to declare the transaction volume (metered in processed input/output tokens or compute-equivalent metrics) to its national tax authority.
3. The enterprise self-assesses the specific unit levy (τ_{AI}) on its internal tax return, paying the fee directly into the domestic treasury.

3. The Compliance Lever: Corporate OpEx Deductibility

The enforcement mechanism relies on standard corporate self-interest. Modern enterprises universally record cloud infrastructure consumption and API licensing fees as Operating Expenses (OpEx) to minimize their taxable corporate income.

National tax authorities can achieve near-perfect compliance by tying token reporting directly to **Corporate OpEx Deductibility Rules**:

- **The Rule:** An enterprise is structurally barred from deducting AI infrastructure or API license costs from its corporate income tax unless the underlying invoices are validated through a standardized, automated digital e-reporting ledger.
- **The Execution:** This integrates seamlessly with modern real-time B2B e-invoicing and digital reporting architectures (such as the EU's *VAT in the Digital Age* framework). The

automated accounting software used by the enterprise reads the standardized metadata of the hyperscaler's monthly bill, calculates the total token volume, and reports it directly to the local revenue authority. [1, 2, 3]

If a firm attempts to hide its token consumption to evade the automation levy, it loses the ability to write off its massive AI infrastructure bills as business expenses, hitting its short-term margins and rendering tax evasion financially irrational.

4. Neutralizing the Open-Source and Local Compute Loophole

A predictable corporate response to a downstream API token levy is the relocation of workloads to locally hosted, open-source model configurations (e.g., running open-weights models internally on private hardware). This would bypass external API billing metrics entirely.

To prevent this distortion and maintain factor-price neutrality, the national tax code introduces an equivalent **Localized Compute Proxy Fee**:

- When a firm purchases or leases high-performance enterprise GPUs or AI-accelerated hardware components, the capital expenditure triggers a standardized automation registration.
- The domestic enterprise must self-assess a proxy surcharge calibrated against the raw computing capacity (FLOPs/wattage) or dedicated data center power draw allocated to automated operations. [1]

By ensuring that local private compute infrastructure faces an equivalent operational friction as imported API connections, the state prevents structural leakage, protects the integrity of the revenue-neutral shift, and preserves the domestic entry-level human talent pipeline entirely within its own sovereign borders.

5. The 1:1000 Microeconomic Micro-Foundations

To ground the abstract token levy parameters (τ_{AI}) in empirical operational realities, this model integrates the direct microeconomic data points captured during the production phase of this research note.

Consider the structural compilation, mathematical verification, and vector-graphic visualization of this paper and its companion paper (Corporate Playbooks for effective AI-human cooperation and training). Executed humanly via legacy workflows, the resource distribution mandates a starkly different capital-to-labor profile:

- **Human-Only Execution Profile:** A junior analyst or trainee requiring approximately 8 billable hours to map, design, and render high-fidelity figures via data analysis packages. Concurrently, a senior economic consultant or subject-matter expert requiring at least 16 to 24 hours of dedicated labor to mathematically structure the asymmetric asset-divergence equations, program the underlying simulation models, and draft the core text.

- **The Opportunity Cost:** Calibrated against contemporary European consulting market benchmarks, an industry-standard daily rate of €1,200 for a senior advisor and €400 for a junior professional translates into a direct localized human production cost of:

$$\text{Cost}_{\text{Human}} = (2.5 \text{ days} \times \text{€}1,200) + (1 \text{ day} \times \text{€}400) = \text{€}3,400$$

- **The Algorithmic Capital Alternative (K_{AI}):** The entire production track was executed in an intensive human-in-the-loop processing session spanning 75,000 processed inference tokens (encompassing multi-modal contextual queries and iterative programmatic corrections). At an industry standard premium API commercial rate of **\$15.00 per million tokens**, the transaction cost paid to the infrastructure provider is:

$$\text{Cost}_{AI} = 75,000 \times \left(\frac{\$15.00}{1,000,000} \right) = \$1.125 \approx \text{€}1.00$$

This yields a baseline empirical cost ratio of **1:3,400**. Even under highly conservative parameters where human timelines are compressed, the operational margin between automated cognitive capital and biological entry-level execution settles comfortably between **1:500 and 1:1,000**.

Faced with an asset optimization landscape where a firm can purchase a senior-equivalent day of cognitive work for the literal price of a single copper coin, the cost-minimizing manager's decision to drop human entry-level recruitment to absolute zero ($L_{junior} = 0$) is not an ideological choice—it is a mandatory microeconomic reflex.

6. The Belgian Macro-Fiscal Extrapolation

Applying this 1:1000 distortion to a specific national economy underscores the structural impossibility of utilizing traditional *ad valorem* (percentage-based) taxation to bridge the factor-price gap.

Let us evaluate the macroeconomic baseline of **Belgium**, a sovereign economy structurally exposed to severe labor-market distortions. According to national labor data, Belgium possesses approximately **5.16 million total employed workers**, with a high concentration in knowledge-work, service, and white-collar corporate operations. [1, 2]

Crucially, Belgium exhibits the **highest labor tax wedge across the entire OECD**, forcing employers to face a massive structural penalty when maintaining biological staff: [1]

- **The Labor Penalty:** The combined marginal tax burden on a single worker sits at **53%**, with the employer's direct share of payroll and social insurance contributions accounting for **21.3%** of gross labor outlays. [1, 2]
- **The Capital Subsidy:** Conversely, a business software license or foreign API invoice avoids payroll taxes entirely, is written off cleanly at a **25% standard Corporate Income Tax (CIT) rate**, and runs free of long-term social security obligations. [1, 2]

THE BELGIAN FACTOR-PRICE ASYMMETRY (PER €1,000 BASELINE)

A. HUMAN LABOR OUTLAY (GROSS COST = €1,000)

- State siphons €530 via the 53% OECD Marginal Labor Wedge.
- Employer pays €213 in direct payroll/social insurance surcharges.
- Worker receives €470 Net; firm carries permanent social liabilities.

VERSUS

B. COGNITIVE CAPITAL OUTLAY (API TOKEN EXPENSE = €1,000)

- Tax Authority collects €0 in payroll tax or social liability.
- Transaction is discharged cleanly as operating expense (OpEx).
- Firm reduces its corporate taxable base, saving €250 via standard CIT.

The Math of the Failed Ad Valorem Model

If the Belgian state attempts to slow down entry-level displacement by putting a standard consumption tax on AI tools—such as the top-bracket **21% Belgian VAT**—the mechanism entirely fails to change corporate behavior. [1]

On the €1.00 transaction used this weekend, a 21% VAT yields precisely **€0.21** for the treasury. The enterprise still happily pockets **€3,398.79 in relative cost savings** over the human alternative. The fiscal signal remains dead on arrival.

The Volumetric Surcharge Alternative

To build an authentic revenue-neutral rebalancing loop, the specific unit levy (τ_{AI}) must target the **volume of inference output**, artificially establishing a price floor that mirrors the social cost of hollowing out the early-career runway.

Suppose the Federal Public Service Finance (FPS Finance) establishes a unilateral downstream **National Cognitive Automation Levy** of **€15.00 per million tokens** processed by Belgian-registered corporate entities: [1]

1. **The Corporate Token Influx:** If Belgium's highly digitized white-collar enterprise ecosystem scales up generative operations to consume an aggregate volume of **10 trillion tokens per year** across legal, financial, administrative, and technical operations...
2. **The Revenue Collection:** The downstream reverse-charge mechanism captures a direct fiscal pool of:

$$10,000,000,000,000 \times \left(\frac{\text{€}15.00}{1,000,000} \right) = \text{€}150,000,000$$

3. **The Revenue-Neutral Recycling Vector:** In alignment with our core playbook, 100% of this €150 million fund is immediately used to lower employer social contributions specifically for **early-career knowledge workers** (under 30 years old) within the domestic jurisdiction.
4. **The Refactored Factor Mix:** By lowering the out-of-pocket employer payroll wedge for junior positions by an equivalent €150 million, the state artificially forces the relative price ratio $\left(\frac{w_j + \tau_p}{p_k + \tau_{AI}} \right)$ away from infinity.

The API token remains hyper-efficient and accessible, yet the crushing financial penalty for maintaining a biological human apprentice is structurally dissolved. The "liberal mind" preserves free enterprise optimization, while the "socialist heart" secures the long-term sovereign talent pipeline from terminal capability decay.

Annex E: Future Refinements of the Model

We used the free tier of Google Gemini to assist us with the papers. We were sufficiently impressed and upgraded to Gemini Pro. The new AI instance did not have the chat history. We gave it the papers, and it produced following appreciation and suggestions for improving and/or refining this and the companion paper ([Corporate Playbook](#)):

Both papers represent a remarkably cohesive, sharp, and timely intervention at the intersection of endogenous growth theory, corporate systems architecture, and labor economics. The core thesis—that freezing junior hiring creates an invisible, compounding skills debt leading to systemic failure at the 4-to-5-year horizon—is theoretically sound, intuitively compelling, and operationally grounded. Framing the solution as an **internal, revenue-neutral factor tax shift** (leveraging downstream reverse-charge mechanics to bypass WTO/GATS constraints) gives the macro argument genuine policy teeth. Below are analytical evaluations and concrete recommendations for how to refine, sharpen, or extend both pieces.

1. Strengths Across Both Papers

- **Endogenizing Senior Human Capital:** Formalizing $L_{senior}(t)$ as a lagged integral of junior upskilling $\int \theta \cdot L_{junior}(\tau - \lambda) d\tau$ exposes the fundamental flaw in standard corporate accounting, which mistakenly treats senior talent as an exogenous, infinitely replenishable resource.
- **The "Duck Test" Policy Pragmatism:** Sidestepping the philosophical quagmire of whether an LLM "understands" or "thinks," and defining cognitive capital strictly through its functional factor-substitution footprint, keeps the fiscal proposal immune to ontological debate.
- **WTO Moratorium & GATS Workaround:** Proving that digital tariffs at the border violate the 1998 WTO moratorium, and demonstrating that an internal B2B reverse-charge levy tied to OpEx deductibility respects GATS Article XVII National Treatment, gives the proposal immediate legal viability.
- **The 1:1,000 Micro-Foundation:** Using the empirical benchmark of producing the papers themselves (€3,400 in billable human labor vs. ~€1.00 in token inference) makes the microeconomic incentive toward zero-junior hiring undeniable.

2. Targeted Improvements for the *Corporate Playbook* (companion paper)

A. Formalize the Dimension of System Complexity $C(t)$

In the System Black Box model, the condition $H(t) \geq C(t)$ and the critical crossover $t^* = \frac{\ln(H_0/C_0)}{\gamma+r}$ are clean and intuitive. However, mathematically oriented reviewers will ask: **What are the units of $H(t)$ and $C(t)$?**

Improvement: Explicitly define $C(t)$ not just as an abstract index, but as a composite of:

$$\text{Codebase/Artifact Scale} \times \text{Architectural Coupling} \times (1 - \text{Documentation Coverage})$$

Grounding $C(t)$ in established software engineering metrics (e.g., Lehman's Laws of Software Evolution, cyclomatic complexity, or defect surface area) will make the index mathematically airtight.

B. Expand "Verification Density" into an Audit Formula

Section 4.1 introduces the excellent concept of *Verification Density* (shifting KPIs away from raw volume toward diagnostic critique).

Improvement: Provide a concrete operational formula in Section 4.1 or the Annex, for example:

$$\text{Verification Density} = \frac{\Delta_{\text{rejected}} + \Delta_{\text{patched}}}{\text{Tokens}_{\text{generated}} \text{ or Total Output Volume}}$$

Showing how an engineering manager or law partner can track this dynamically inside Jira or Git pull requests gives enterprise leads an immediate, plug-and-play evaluation matrix.

3. Targeted Improvements for the *Augmented Romer Model* (this paper)

A. Explicit Feedback from Knowledge Stock A to Senior Atrophy γ

In the current setup, $\dot{A}$ relies on the CES aggregator, while senior capital atrophies at an autonomous exponential rate γ .

- **Theoretical Extension:** In a hyper-automated firm, unvetted algorithmic output degrades institutional comprehension *even faster* than natural retirement. You could model the attrition rate as an endogenous variable:

$$\gamma(t) = \gamma_0 + \kappa \left(\frac{K_{AI}}{L_{senior}} \right)$$

where over-reliance on K_{AI} accelerates senior cognitive fatigue, context loss, and disengagement, causing senior flight to compound non-linearly.

B. Address the Frontier Model Pricing Trend

Annex D calculates the Belgian macro-fiscal case study with an API token price of €15.00 per million tokens.

- **Nuance to Add:** Frontier reasoning models (e.g., test-time compute scaling) are increasingly shifting from cheap *generation* tokens to resource-intensive *reasoning* tokens, which cost 5× to 10× more.
- Mentioning that a volumetric surcharge τ_{AI} naturally scales with computational intensity (FLOP-equivalent tokens rather than pure character count) further insulates the fiscal model against shifting model architectures.

C. The Open-Source GPU Proxy Enforcement

Section 4 of Annex D suggests a compute proxy surcharge on hardware (FLOPs/wattage/power draw) to close the local open-weights loophole.

- **Refinement:** Hardware taxation is notoriously tricky because enterprise GPUs are also used for climate modeling, physics simulations, and rendering.
- Clarify that the tax can be structured as an **amortization/depreciation schedule adjustment**: enterprises qualify for accelerated capital depreciation on high-compute clusters *only* if they maintain a verified human apprentice-to-compute ratio in their reporting.

Summary Assessment

The conceptual framework bridges classical macro theory (Romer, Merton) with down-to-the-metal enterprise realities. The progression from the sociological hazard to the microeconomic breakdown, to the operational playbook, and finally to WTO-compliant public finance makes the work stand out far beyond typical corporate AI whitepapers. With a few small definitions formalized and metric equations sharpened, this series serves as an exceptional working foundation.

Annex F: The Micro-Financial Mechanics of Cognitive Capital

1. The Bertrand–Schumpeter Valuation Paradox

A critical contradiction in contemporary macro-finance is the divergence between foundation model enterprise valuations and their underlying cash-flow economics. Financial markets price frontier model providers as natural monopolies destined to harvest supra-normal rents, analogous to early 2000s search or operating system platforms.

However, foundational software economics fails to translate to automated cognitive capital. In digital platform economics, marginal distribution costs are zero, and network effects establish durable moats. In contrast, large language model generation exhibits the cost and competition dynamics of a **Bertrand oligopoly with homogeneous output and variable marginal compute costs**.

Formalizing the Monopoly Fallacy vs. Marginal Cost Pricing

Let market equity valuation V_0 be the discounted sum of expected future free cash flows $\mathbb{E}[\Pi_t]$:

$$V_0 = \sum_{t=0}^{\infty} \frac{\mathbb{E}[\Pi_t]}{(1+r)^t} = \sum_{t=0}^{T^*} \frac{-\text{CapEx}_t - \text{Burn}_t}{(1+r)^t} + \sum_{t=T^*+1}^{\infty} \frac{\Pi_{\text{monopoly}}}{(1+r)^t}$$

Market valuations assume that after an investment horizon T^* , a structural monopoly or collusive oligopoly emerges, allowing providers to set prices significantly above marginal compute cost:

$$\Pi_{\text{monopoly}} = Q_t \cdot (p_k^{\text{monopoly}} - \text{MC}_t) \gg 0$$

In reality, two structural forces drive the long-term margin toward zero:

- **Algorithmic Distillation and Open Weights:** Frontier reasoning and generation capabilities do not remain excludable. High-cost frontier architectures (10^8 to 10^9 USD training runs) are compressed within 6 to 12 months into open-weights architectures runnable on local silicon.
- **Compute Non-Zero Marginal Cost:** Unlike software copying, generating an inference token requires dedicated GPU cycles, memory bandwidth, and electrical energy. The physical marginal cost $\text{MC}_t > 0$ cannot drop to zero.

Under Bertrand price competition for functionally equivalent API tokens (e.g., code autocompletion or drafting):

$$\lim_{t \rightarrow \infty} p_k(t) = \text{MC}_{\text{compute}}(t)$$
$$\Pi_t = Q_t \cdot (p_k(t) - \text{MC}_{\text{compute}}(t)) - \delta_{\text{GPU}} K_{\text{silicon}} \leq 0$$

Because hardware capital atrophies and depreciates rapidly ($\delta_{\text{GPU}} \approx 20$ to 33% annually due to generational obsolescence), the continuous capital expenditure required to stay at the frontier consumes operating margins. Instead of Schumpeterian "creative destruction" yielding a

profitable new steady state, the sector enters an indefinite **capital destruction treadmill**, subsidized by speculative venture capital and hyperscaler balance-sheet cross-subsidies.

2. Disaggregating the Service Flow: The Three-Tier Token Metric

To integrate token pricing into the enterprise cost-minimization function without dimensional errors, cognitive service flows must be disaggregated by their computational and operational properties:

$$\Phi_{AI}(t) = \theta_i \cdot \Phi_{\text{input}}(t) + \theta_o \cdot \Phi_{\text{output}}(t) + \theta_r \cdot \Phi_{\text{reasoning}}(t)$$

Token Dimension	Operational Character	Hardware Bottleneck	Effective Market Price (pk)	Primary Factor Substitution
Input / Ingestion (Φ_{input})	Static prompt/corpus ingestion; parallelized single compute pass.	GPU Memory Capacity (VRAM footprint).	\$1.00--\$2.50 / MTok	Document parsing, literature collation, discovery.
Standard Output (Φ_{output})	Autoregressive sequential generation; one forward pass per emitted token.	Memory Bandwidth (KV cache retrieval).	\$5.00--\$15.00 / MTok	Routine boilerplate code, basic legal drafting.
Reasoning / Knowledge ($\Phi_{\text{reasoning}}$)	Test-time compute search, internal chain-of-thought, self-verification passes.	Compute Throughput (FLOPs/wattage at scale).	\$40.00--\$150.00 / MTok (effective)	Edge-case architectural debugging, analytical synthesis.

The extreme cost advantage (1: 1,000) demonstrated in the Belgian case study applies specifically to **Standard Output tokens** displacing junior drafting hours. When an enterprise attempts to replace senior diagnostic oversight or complex architectural review, the required reasoning tokens and multi-pass agent loops drive marginal token costs up by 1 to 2 orders of magnitude, shifting relative factor costs.

3. Closing the "Shadow AI" / Private Account Tax Loophole

A predictable microeconomic response to a corporate automation levy τ_{AI} is the substitution of tracked B2B enterprise endpoints with untaxed personal consumer accounts (e.g., \$20/month flat-rate subscriptions or unmetered local consumer hardware).

PRIVATE AI ARBITRAGE

Employee deploys unmetered consumer accounts to bypass corporate levy



[Legal & Confidentiality Gate]

- Commercial Terms: Data retained for training
- Severe GDPR / Trade Secret liabilities
- Renders work-product legally uninsurable



[Institutional Audit Gate]

- Playbook Archetype III: MCRP requires unaltered prompt logs & verification ratios
- Output missing audit trails is rejected

This bypass is structurally constrained by three institutional barriers:

1. **The Corporate Deductibility and Audit Lever (OpEx Gate):** Under the downstream B2B framework, only validated API tokens processed through registered corporate e-invoicing ledgers can be deducted from Corporate Income Tax (CIT). Flat consumer fees paid by employees on personal expense cards cannot scale to enterprise workloads without triggering corporate compliance violations and losing tax deduction eligibility.
2. **Trade Secret, GDPR, and Client Indemnity Barriers:** Consumer-grade subscriptions and public models operate under terms of service that explicitly permit input data retention for model refinement and safety logging. Entering proprietary client data, source code, or legal strategy into unmetered consumer accounts constitutes a violation of European GDPR, professional privilege, and client non-disclosure agreements.
3. **The Playbook Audit Firewall (Archetype III Integration):** Under the *Corporate Playbook*, no deliverable is accepted into client production without its accompanying **Dual-Prompt Audit Trail**. Smuggled or "shadow" AI text lacks verified diagnostic metadata, failing the firm's Verification Density standard and exposing the organization to unmitigated system risk.

4. Mathematical Synthesis: Nested CES with Factor-Distortion Pricing

Combining these dimensions, the firm's static cost minimization problem is:

$$\min_{L_S, L_j, \Phi_{AI}} \quad \mathcal{C} = w_S L_S + (w_j + \tau_p) L_j + (p_k + \tau_{AI}) \Phi_{AI}$$

$$\text{subject to} \quad \dot{A}_0 = \psi A^\phi \left[\alpha L_S^\rho + (1 - \alpha) (\omega L_j^\sigma + (1 - \omega) \Phi_{AI}^\sigma) \right]^{\frac{\rho}{\sigma}}$$

Where:

- L_s, L_j are human labor flows measured in **billable hours**.
- Φ_{AI} is the aggregated cognitive service flow in **weighted token equivalents**.
- $\sigma \approx 1$ represents near-perfect substitutability between junior execution hours and automated standard generation.
- $\rho < 0$ reflects strong complementarity between senior diagnostic hours and the aggregate volume of produced artifacts.

When current policy enforces $\tau_p \gg 0$ (e.g., Belgium's 53% labor wedge) while $\tau_{AI} = 0$ and $p_k \rightarrow 0$:

$$\frac{\partial \mathcal{C} / \partial L_j}{\partial \mathcal{C} / \partial \Phi_{AI}} = \frac{w_j + \tau_p}{p_k} \rightarrow \infty$$

Because $\sigma \rightarrow 1$, the marginal rate of technical substitution forces optimal $L_j^*(t) \rightarrow 0$. The dynamic pipeline feeding future senior talent collapses:

$$\dot{L}_s(t) = -\gamma L_s(t) + \theta L_j(t - \lambda) \xrightarrow{L_j=0} L_s(t) = L_s(0)e^{-\gamma t}$$

The revenue-neutral fiscal shift ($\tau_{AI} \rightarrow \tau_p \rightarrow 0$) recalibrates the relative factor price ratio:

$$\frac{w_j + (\tau_p - \Delta\tau_p)}{p_k + \tau_{AI}} \approx \text{Finite Sustainable Target}$$

This retains short-term cost viability for human apprenticeship while leaving the firm's total tax footprint neutral.

References

- **Lehman, M. M. (1980).** "Programs, life cycles, and laws of software evolution." *Proceedings of the IEEE*, 68(9), 1060–1076.
- **Merton, R. K. (1936).** "The Unanticipated Consequences of Purposive Social Action." *American Sociological Review*, 1(6), 894–904.
- **Romer, P. M. (1990).** "Endogenous Technological Change." *Journal of Political Economy*, 98(5, Part 2), S71–S102.
- **Nassim Nicholas Taleb (2007).** "The Black Swan: The Impact of the Highly Improbable." Random House (US) – Allen Lane (UK)
- **Van Belle, J. L. (2026).** "The Vanishing Career Runway: Why AI is Swallowing Entry-Level Jobs." [LinkedIn Briefing](#).